

## *Contents of Lecture 2*

- The Scala programming language
- Actors
- Parallel programming with Akka

# Purpose of this lecture

- That you will understand why Scala may be interesting
- You will understand the key concepts of message passing using **actors**, introduced by Carl Hewitt at MIT 1973.
- You will understand enough about Scala actors that you can write a parallel version of the preflow-push algorithm in Scala

# Scala and Akka

- Martin Odersky designed Generic Java and the Java compiler `javac` for Sun. He is a professor at EPFL in Lausanne.
- The Scala language produces Java byte code and Scala programs can use existing Java classes.
- When you run a Scala program, the JVM cannot see any difference between Java and Scala code.
- A good source to start with Akka:  
`https://developer.lightbend.com/start`
- Or start with the example at Tresorit
- Akka was created by Jonas Bonér from Sweden
- Lightbend is a company founded by Odersky, Bonér and another person.
- Download sbt: `https://www.scala-sbt.org/download.html`

# Scala is a functional and object oriented language

- It is intended to be **scalable** and suitable to use from very small to very large programs.
- The Scala compiler, `scalac`, usually can infer the types of variables so you don't have to type them.

```
var capital = Map("Denmark" -> "Copenhagen", "France" -> "Paris", "Sweden" -> "Stockholm");  
capital += ("Germany" -> "Berlin");  
println(capital("Sweden"));
```

- There is no need to declare the type of the variable `capital` since `scalac` can do it for you.
- Less typing can potentially lead to faster programming — at least if the tedious part of the typing can be eliminated.

# Shorter class declarations

- A short class declaration in Scala:

```
class Example(index: Int, name: String)
```

- The same class in Java:

```
class Example {  
    private int index;  
    private String name;  
  
    public Example(int index, String name) {  
        this.index = index;  
        this.name = name;  
    }  
}
```

# Redefinable operators as in C++

```
def factorial(x: BigInt): BigInt = if (x == 0) 1 else x * factorial(x - 1)
```

- A function is defined using `def` and `=` and an expression.
- Or using the Java class `BigInteger`:

```
import java.math.BigInteger;

def factorial(x: BigInteger): BigInteger =
  if (x == BigInteger.ZERO)
    BigInteger.ONE
  else
    x.multiply(factorial(x.subtract(BigInteger.ONE)))
```

- It is obvious which is nicer.

# Scala is statically typed

- Lisp, Smalltalk, Ruby, Python and many other languages are dynamically typed, which means type checking is performed at runtime.
- Scala and to a very large extent also C are statically typed.
- Of course, C is a very small language and much easier to type check.
- For C, if you use `<stdarg.h>` (which you usually shouldn't) or insane casts (which result in undefined behaviour = serious bug) the C compiler will not help you.
- Look at this program:

```
(defun sumlist (h)
  (if (null h) 0 (+ (car h) (sumlist (cdr h)))))

(setq b '(1 2 3 4))
(setq c '("x" "y" "z"))

(print (sumlist b))
(print (sumlist c))
```

- Which language is it?
- When is the error detected during dynamic type checking?

# Answers

```
(defun sumlist (h)
  (if (null h) 0 (+ (car h) (sumlist (cdr h)))))
```

```
(setq b '(1 2 3 4))
(setq c '("x" "y" "z"))
```

```
(print (sumlist b))
(print (sumlist c))
```

- The language is Common Lisp.
- The error is detected when adding the string "z" to zero: Old measurement from several years back (also for C below)

```
> time clisp a.lisp
```

```
10
```

```
*** - +: "z" is not a number
```

```
real      0m0.062s
user      0m0.045s
sys       0m0.017s
```



# A C Compiler must issue a diagnostic message

```
#include <stdlib.h>

typedef struct list_t  list_t;

struct list_t {
    list_t*    next;
    int        value;
};

list_t* cons(int value, list_t* list)
{
    list_t*    p;

    p = malloc(sizeof(list_t));
    if (p == NULL)
        abort();
    p->value = value;
    p->next = list;

    return p;
}

int sumlist(list_t* h)
{
    return h == NULL ? 0 : h->value + sumlist(h->next);
}

int main(void)
{
    list_t*    p;
    list_t*    q;

    p = cons(1, cons(2, cons(3, cons(4, NULL))));
    q = cons("x", cons("y", cons("z", NULL)));    // static type error
}
```

# GCC output

```
> time gcc a.c
a.c: In function 'main':
a.c:35: warning: passing argument 1 of 'cons' makes integer from pointer without a cast
a.c:11: note: expected 'int' but argument is of type 'char *'
a.c:35: warning: passing argument 1 of 'cons' makes integer from pointer without a cast
a.c:11: note: expected 'int' but argument is of type 'char *'
a.c:35: warning: passing argument 1 of 'cons' makes integer from pointer without a cast
a.c:11: note: expected 'int' but argument is of type 'char *'

real    0m0.150s
user    0m0.103s
sys     0m0.047s
```

# Clang output

```
> time clang -S a.c
a.c:35:31: warning: incompatible pointer to integer conversion passing
      'char [2]' to parameter of type 'int'
      q = cons("x", cons("y", cons("z", NULL)));
                                ^~~~

a.c:11:18: note: passing argument to parameter 'value' here
list_t* cons(int value, list_t* list)
             ^

a.c:35:21: warning: incompatible pointer to integer conversion passing
      'char [2]' to parameter of type 'int'
      q = cons("x", cons("y", cons("z", NULL)));
                        ^~~~

a.c:11:18: note: passing argument to parameter 'value' here
list_t* cons(int value, list_t* list)
             ^

a.c:35:11: warning: incompatible pointer to integer conversion passing
      'char [2]' to parameter of type 'int'
      q = cons("x", cons("y", cons("z", NULL)));
              ^~~~

a.c:11:18: note: passing argument to parameter 'value' here
list_t* cons(int value, list_t* list)
             ^

3 warnings generated.

real    0m0.050s
user    0m0.030s
sys     0m0.020s
```

# A Scala compiler must issue a diagnostic message

```
class Test {  
  def sumlist(h:List[Int]) : Int = if (h.isEmpty) 0 else h.head + sumlist(h.tail);  
  var a = List(1,2,3,4);  
  var b = sumlist(a);  
  var c = List("x", "y", "z");  
  var d = sumlist(c);  
}  
  
> time scalac a.scala  
a.scala:6: error: type mismatch;  
found    : List[java.lang.String]  
required: List[Int]  
  var d = sumlist(c);  
                ^  
  
one error found  
  
real    0m1.697s  
user    0m4.384s  
sys     0m0.088s
```

- Measurement made on login.student.lth.se September 12, 2019.
- The type analysis for Scala is of course more complex than for C.
- Avoid compilers with this compilation speed for large source code
- Compilation speed of Scala may improve significantly in the future.
- Previous measurement was 11 s.

# The Fast Scala compiler

- There is a server program `fsc` which is faster than `scalac` because it avoids some initializations.
- Clang and Common Lisp were fastest.
- `clisp` was originally written in assembler and Lisp for Atari machines but has been rewritten in portable C and Lisp.

# The Scala build tool: sbt

- sbt is a tool which downloads required libraries, starts the Scala compiler and runs the program
- It will try to make a Scala program by compiling all Scala files in the current directory — so keep only one version of your program there!
- In the lab it is sufficient to type make. Abbreviated output:

```
$ make  
./sbt run < i  
[info] loading settings for project lab1 from build.sbt  
[info] running main  
f = 9924  
t = 3.85 s
```

- It takes about 10 s to compile the 474 lines of Scala code

# Scala basics: val vs var

- Writing `var a = 1`, we declare an initialized `Int` variable that we can modify.
- With `val a = 1`, `a` becomes readonly instead.
- The following declares an array:

```
val a = new Array[String](2);  
a(0)  = "hello";  
a(1)  = "there";
```

- Note that it is `a` that is readonly, not its elements.
- We can iterate through an array like this, for example:

```
for (s <- a) println(s);
```

- We should not declare the variable `s`.

# Numbers are objects

- Consider

```
for (i <- 0 to 9) println(i);
```

- Here the zero actually is an object with the method `to`.
- In many cases a method name can be written without the dot but rather as an operator.



# Companion classes

- In Java you can have static attributes of a class which are shared by all objects of that class.
- In Scala, you instead create a **companion class** with the keyword **object** instead of **class**:

```
object A {  
    var a = 44;  
}  
  
class A {  
    println("a = " + A.a);  
}  
  
object Main {  
    def main(args: Array[String]) {  
        val a = new A;  
    }  
}
```

- By default attributes are public in all Scala classes, but an object may access a private attribute of its companion class.

# Class declarations with attributes

- You can declare a class like this:

```
class B(u: Int, var v: Int) {  
    def f = u + v;  
}
```

- The parameters of a constructor by default become `val` attributes of the class.
- Therefore only `v` can be modified.
- Even if you only need the parameters in the constructor, they become attributes and cheerfully consume memory for you.

# Code reuse in Scala using traits

- Scala uses single inheritance as Java with the same keyword `extends`.
- Instead of Java's interfaces which only provide abstract methods, Scala has the concept of a trait.
- Unlike an interface, a trait can contain attributes and code, however.
- A trait is similar to a class except that the constructor cannot have parameters.

```
object Main {  
  def main(args: Array[String]) {  
  
    val a = new C(44);  
  
    a.hello;  
    a.bye;  
  }  
}  
  
class A(u: Int) {  
  def bye { println("bye bye with u = " + u); }  
}  
  
trait B {  
  def hi { println("hello"); }  
}  
  
class C(v: Int) extends A(v) with B {  
  def hello { hi; }  
}
```

- The standard class `List` is singly linked and consists of a pair of data and a pointer to the next element.
- An empty list is written either as `Nil` or `List()`.
- Five (of many) methods are:
  - `::` — create a list: `val h = 1 :: 2 :: 3 :: Nil`, which means:  
`val h = (1 :: (2 :: (3 :: Nil)))`.  
This can also be written as `val h = List(1, 2, 3)`
  - `:::` — create a new list by concatenating two lists.  

```
val a = List(1, 2, 3);  
val b = List(4, 5, 6);  
val c = a ::: b;
```
  - `isEmpty` — boolean
  - `head` — data in first element
  - `tail` — the rest of the list starting with the 2nd element.

# Reversing a List 1(2)

```
def rev[T](h:List[T]) : List[T] = {  
  if (h.isEmpty)  
    h;  
  else  
    rev(h.tail) ::: List(h.head);  
}
```

- This function is generic with element type T.
- It's not the most efficient way to reverse a list since list concatenation must traverse the left operand list.
- This version of reverse has quadratic time complexity.
- How can we do it in linear time?

## Reversing a List 2(2)

```
def rev1[T](h : List[T], q : List[T]) : List[T] = {  
  if (h.isEmpty)  
    q;  
  else  
    rev1(h.tail, h.head :: q);  
}
```

```
def rev[T](h : List[T]) : List[T] = {  
  rev1(h, List());  
}
```

- Better.

# Pattern matching in Scala

- Pattern matching means we provide a sequence of cases against which the input data is matched.
- The first case that matches the input data is executed.

```
def rev[T](xs: List[T]) : List[T] = xs match {  
  case List()    => xs;  
  case x :: xs1 => rev(xs1) :: List(x);  
}
```

- The reverse of the empty list is the parameter `xs`.
- The non-empty list matches a list with at least one element, as in the second case.
- Pattern matching is used extensively in functional programming.
- We will use pattern matching when receiving actor messages.

# Programming with actors in Scala

- Programming with actors is in one sense just writing another multithreaded program.
- In another sense it's completely different because you should use no locks or condition variables or the like.
- An actor is like a thread which sits and waits for a message to arrive.



# Sending a message

- With actors, messages are sent to an actor — and not to a mailbox or channel as in other systems.
- A message can be a variable or a type
- Assume an actor A has a reference to another actor B, then A can send messages of types C and D to B using:

```
val e = 124
B ! C;
B ! D(42);
B ! e;
```

- The sending actor immediately continues execution without waiting for the message to arrive.
- Without a parameter, the message type should be declared as:  
case object C
- With a parameter, instead use: case class D(x:Int)

# Receiving a message

```
case object C
case class D(x: Int)
case object Thanks

def receive = {

  case D(x: Int) => println("got a D with x = " + x)

  case C          => { println("got a C"); sender ! Thanks }

  case e:Int      => println("got an Int e = " + e)

}
```

- The sending actor is sender
- Case classes and objects must start with a capital otherwise they become variables which match everything
- When there are problems, use a variable in a last case and print it

# Evaluating a message

- With pattern matching on the message, the action to perform is selected.
- If there is no match, the message is discarded.
- After performing it, the actor repeats the waiting for another message,
- It's not necessarily easier to program with actors than with locks.
- For instance, you can end up with a deadlock if two actors are waiting for messages from each other.
- A message never interrupts an actor — the actor processes a message to completion before processing the next message

# Message arrival

- In some actor-based systems the arrival order of messages is not specified — even for messages from the same sender.
- For Scala actors, messages sent from one actor to another always arrive in the same order as they were sent.
- When an actor has created/modified data and sent a message to another actor, that data will be visible through the cache memory in the receiving thread so there is no need to use volatile as in Java
- As we will see, there is a volatile keyword in C/C++ as well but it means something else

# Sharing mutable data and data-races

- If you share a message with data that both actors may want to modify, you can have a data-race
- For preflow push, for an edge  $(u, v)$  both nodes may at some point want to modify the flow of their edge
- If in your implementation both  $u$  and  $v$  can modify the flow at the same point in time, you have a data-race
- The nodes should be actors but although the edges could be actors as well, the program will be unnecessarily complicated then so use normal objects for the edges

# Sharing data between actors in general

- This is not so much relevant for the lab but still very important
- Suppose actors have vectors or other larger data which they update and sometimes share with other actors
- When an actor  $X$  asks for the vector of another actor  $Y$  and:
  - $Y$  is fine with giving the vector to  $X$ , but
  - $Y$  wants to continue modifying the vector, then
- instead of creating a data-race by having both  $X$  and  $Y$  access the vector we can do as follows:
  - $Y$  can make a copy of the vector and give it to  $X$  so  $X$  can do what it wants
- If  $X$  and possibly other actors only want to read the vector we can instead do:
  - $Y$  can make a copy of the vector that it shares with any actor as immutable data (immutable = readonly)
  - $Y$  can have one internal vector that it modifies itself and when sufficiently done it can copy it and let future actor requests see that instead

# Declaring and creating an Actor

- An actor class can be declared as:

```
class Node(val index: Int) extends Actor {  
    var      e = 0;  // excess preflow  
    var      h = 0;  // height  
}
```

- But we do not make an array  $n$  of Node objects:

```
var      node: Array[ActorRef] = null  
node = new Array[ActorRef](n)  
  
for (i <- 0 to n-1)  
    node(i) = system.actorOf(Props(new Node(i)), name = "v" + i)
```

- So we create an array of ActorRef
- A "factory" creates each actor, i.e. not simply new Node(i)
- We will soon see what system is

# A Preflow actor

- It is often convenient to have a central controller which determines when an algorithm is finished.
- Two disadvantages with this approach are
  - ① A single controller can be a performance bottleneck — risk in lab 1
  - ② A single point of failure can cause a service to fail — ignore in lab 1
- Usually when an algorithm is parallelized to be run on a multicore computer we assume crashes are software bugs, i.e., none in our code
- In a distributed system we need to be more careful
- Such care obviously creates overhead which we want to avoid in a multicore
- When only the sink has a positive excess preflow, we can stop



# The Ceberg termination criterion

- The source starts with a negative excess preflow after it has pushed to each neighbor
- Then the source possibly gets some flow back
- The excess preflow of the source increases monotonically
- That is  $|e_f(s)|$  decreases monotonically
- The sink gets more and more, ie  $e_f(t)$  increases monotonically
- When the  $|e_f(s)| = e_f(t)$  no other node can have any excess preflow
- Then we are finished
- This is easier to detect than counting the number of nodes with excess preflow
- Invented by Nils Ceberg (who took the course 2021).
- Although this may seem obvious, current published articles on distributed preflow push don't do this

# Declaring and creating the controller actor

- The controller can be declared without parameters as:

```
class Preflow extends Actor
{
    var s = 0;
    var t = 0;
    var n = 0;
    var edge:Array[Edge] = null
    var node:Array[ActorRef] = null
}
```

- Let the sink tell the source when it has received more preflow

```
val system = ActorSystem("Main")
val control = system.actorOf(Props[Preflow], name = "control")
```

- Without parameters, it is simpler to create an actor
- Note the different syntax compared to creating the node actors

# self vs this

- Suppose we have

```
class A extends Actor { }
```

- Then `this` refers to `A` and `self` to `ActorRef`
- For the controller to send a reference to itself to a node `u`, use `self`:

```
u ! Control(self) // u is an ActorRef in the array node
```

- For a node to save the controller parameter, use `this`:

```
case Control(control:ActorRef) => this.control = control
```

# Waiting for an answer

- We can create a timeout and use ? when sending a message:

```
implicit val t = Timeout(4 seconds);  
val flow = control ? Maxflow  
val f = Await.result(flow, t.duration)  
println("f = " + f)
```

- The type of the sender is ActorRef

```
var      ret:ActorRef          = null  
case Maxflow => {  
    ret = sender                // save sender for a future reply  
    node(s) ! Source(n)        // tell s it is source and has h = n  
    node(t) ! Sink              // tell t it is sink  
    node(s) ! Start             // tell s to do initial pushes  
}
```

# Stopping the actors

- The stop method can be used:

```
system.stop(control);  
for (i <- 0 to n-1)  
    system.stop(node(i))  
system.terminate()
```

# Useful functions of a Node

```
def id: String = "@" + index; // easy to grep or search for
```

```
def status: Unit = {  
    if (debug) println(id + " e = " + e + ", h = " + h);  
}
```

```
def enter(func: String): Unit = {  
    if (debug) { println(id + " enters " + func); status }  
}
```

```
def exit(func: String): Unit = {  
    if (debug) { println(id + " exits " + func); status }  
}
```

```
def relabel : Unit = {  
    enter("relabel")  
    h += 1  
    exit("relabel")  
}
```

# Increasing parallelism when doing push

- Think through what you need to do for a push
- Do you need to wait for a reply?
- What should the reply be in that case?
- If you want to wait for a reply, can you then do multiple pushes concurrently?
- By concurrently here is meant to have sent multiple "push" messages before waiting for a reply from each
- See pdf for lab 1 for requirements to pass the lab
- If this would not be about network flow but instead chat messages to a group, you probably would want to send multiple messages concurrently

# The adjacency list of a Node

```
var      edge: List[Edge] = Nil      // at Node construction

def discharge: Unit = {              // do pushes or relabel
  var      p: List[Edge]    = Nil    // pointer into edge
  var      a: Edge          = null   // edge to work with

  p = edge
  while (p != Nil) {
    a = p.head
    p = p.tail
  }
}
```

- This is one option to iterate through the adjacency list
- Note that you *may* want to wait for a reply before doing the next edge
- Try to increase parallelism by doing multiple pushes before waiting for a reply!



# What to do in the lab

- Take the smallest input, draw the graph on paper, and write down all messages that should be sent.
- When you are happy with that, you may want to try bigger input or start thinking about source code.
- You will get an incomplete program which:
  - measures execution time
  - reads the input and creates the graph
  - asks the controller to compute the maximum flow
  - stops the actors
- You need to write code for:
  - iterate through the adjacency list and do push
  - if no push could be done due to heights of neighbors, do a relabel
  - tell controller when the sink has received more excess preflow

## Hint: actor decision making

- When two or more actors are involved in a decision (such as whether it is the right time to do a push) it is important to figure out who can make the final decision so that no algorithm invariants are violated
- Assume a node  $u$  asks a neighbor  $v$  for its height,
- $v$  replies with a height lower than the height of  $u$ , and
- $u$  pushes to  $v$  but before the push arrives,  $v$  has done a relabel (so the push should not have been done).
- How can you avoid this problem?

## Hint: debugging

- Let each node have a debug attribute that you can enable/disable easily without editing more than one line of source code
- Start with the smallest inputs and compare the output from your program with that from the sequential C program.
- Use grep to find out what a certain node is doing:

```
forsete> cat i0
3 2 0 0
0 1 10
1 2 2
forsete> sbt run < i0 > x
forsete> grep @1 x
```

- This should print out all lines from x in which node 1 does something since node 1 is identified with @1